

```
In [16]: # Import packages
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.cross_decomposition import PLSRegression
from sklearn.ensemble import RandomForestRegressor
from xgboost import XGBRegressor
from sklearn.tree import DecisionTreeRegressor
from sklearn.model_selection import train_test_split, cross_val_score, cross_val_predict, KFold, GridSearchCV,
from sklearn.metrics import mean_squared_error, r2_score
from sklearn.preprocessing import OneHotEncoder, StandardScaler
from scipy.optimize import minimize
import statsmodels.api as sm
import scipy.stats as stats
from sklearn.linear_model import LassoCV
from collections import Counter
from scipy.stats import randint, uniform
```

```
In [16]: !jupyter nbconvert 5291_Project.ipynb --to html

[NbConvertApp] Converting notebook 5291_Project.ipynb to html
[NbConvertApp] Writing 1772301 bytes to 5291_Project.html
```

```
In [3]: import os
curr_path = os.getcwd()
citibike_path = os.path.join(curr_path, "citibike_day.csv")
weather_path = os.path.join(curr_path, "nyc_Weather_2013_2023.csv")
```

```
In [4]: # Load data
citibike_df = pd.read_csv(citibike_path)
weather_df = pd.read_csv(weather_path)
```

```
In [6]: # Display basic information
citibike_df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 365 entries, 0 to 364
Data columns (total 5 columns):
#   Column          Non-Null Count  Dtype
---  ---
0   Unnamed: 0      365 non-null   int64
1   date            365 non-null   object
2   num_trips       365 non-null   int64
3   avg_duration    365 non-null   float64
4   total_duration  365 non-null   int64
dtypes: float64(1), int64(3), object(1)
memory usage: 14.4+ KB
```

```
In [7]: weather_df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4017 entries, 0 to 4016
Data columns (total 15 columns):
#   Column                                Non-Null Count  Dtype
---  ---
0   time                                  4017 non-null   object
1   temperature_2m_max                    4017 non-null   float64
2   temperature_2m_min                    4017 non-null   float64
3   apparent_temperature_max              4017 non-null   float64
4   apparent_temperature_min              4017 non-null   float64
5   precipitation_sum                     4017 non-null   float64
6   rain_sum                             4017 non-null   float64
7   snowfall_sum                         4017 non-null   float64
8   precipitation_hours                   4017 non-null   float64
9   sunshine_duration                    4017 non-null   float64
10  windspeed_10m_max                     4017 non-null   float64
11  windgusts_10m_max                     4017 non-null   float64
12  winddirection_10m_dominant             4017 non-null   int64
13  shortwave_radiation_sum               4017 non-null   float64
14  et0_fao_evapotranspiration            4017 non-null   float64
dtypes: float64(13), int64(1), object(1)
memory usage: 470.9+ KB
```

```
In [5]: # Compute average daily temperature
weather_df['temperature_avg'] = (weather_df['temperature_2m_max'] + weather_df['temperature_2m_min']) / 2
```

```
In [6]: # Rename 'time' to 'date' and convert to datetime
weather_df = weather_df.rename(columns={'time': 'date'})
weather_df['date'] = pd.to_datetime(weather_df['date'])
citibike_df['date'] = pd.to_datetime(citibike_df['date'])
```

```
In [10]: weather_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4017 entries, 0 to 4016
Data columns (total 16 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   date                                  4017 non-null   datetime64[ns]
1   temperature_2m_max                   4017 non-null   float64
2   temperature_2m_min                   4017 non-null   float64
3   apparent_temperature_max             4017 non-null   float64
4   apparent_temperature_min             4017 non-null   float64
5   precipitation_sum                    4017 non-null   float64
6   rain_sum                             4017 non-null   float64
7   snowfall_sum                        4017 non-null   float64
8   precipitation_hours                  4017 non-null   float64
9   sunshine_duration                    4017 non-null   float64
10  windspeed_10m_max                   4017 non-null   float64
11  windgusts_10m_max                   4017 non-null   float64
12  winddirection_10m_dominant          4017 non-null   int64
13  shortwave_radiation_sum             4017 non-null   float64
14  et0_fao_evapotranspiration          4017 non-null   float64
15  temperature_avg                     4017 non-null   float64
dtypes: datetime64[ns](1), float64(14), int64(1)
memory usage: 502.3 KB
```

```
In [7]: # Select and rename relevant columns
weather_selected = weather_df[[
    'date',
    'temperature_avg',
    'precipitation_sum',
    'windspeed_10m_max',
    'sunshine_duration'
]]
```

```
In [8]: # Merge with Citi Bike data
merged_df = pd.merge(citibike_df, weather_selected, on='date', how='inner').dropna()
```

```
In [9]: # Display first few rows of the merged dataset
merged_df.head()
```

```
Out[9]:
```

	Unnamed: 0	date	num_trips	avg_duration	total_duration	temperature_avg	precipitation_sum	windspeed_10m_max	sunshine_duration
0	0	2015-01-01	5317	801.806658	4263206	-1.00	0.0	21.3	29236.59
1	1	2015-01-02	11304	731.240977	8265948	1.45	0.0	19.9	29150.96
2	2	2015-01-03	4478	655.284279	2934363	1.75	16.5	15.8	6444.89
3	3	2015-01-04	7849	679.554211	5333821	10.00	7.2	22.3	0.00
4	4	2015-01-05	14506	637.811802	9252098	1.55	0.0	26.8	29161.29

```
In [14]: merged_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 365 entries, 0 to 364
Data columns (total 9 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Unnamed: 0                            365 non-null   int64
1   date                                  365 non-null   datetime64[ns]
2   num_trips                             365 non-null   int64
3   avg_duration                          365 non-null   float64
4   total_duration                       365 non-null   int64
5   temperature_avg                      365 non-null   float64
6   precipitation_sum                    365 non-null   float64
7   windspeed_10m_max                   365 non-null   float64
8   sunshine_duration                   365 non-null   float64
dtypes: datetime64[ns](1), float64(5), int64(3)
memory usage: 25.8 KB
```

Regression Analysis

```
In [15]: # Define response and predictors
X = merged_df[['temperature_avg', 'precipitation_sum', 'windspeed_10m_max', 'sunshine_duration']]
y = merged_df['num_trips']
```

```
In [16]: # Add constant to predictors
X = sm.add_constant(X)

# Fit the regression model
```

```
model = sm.OLS(y, X).fit()
```

```
In [17]: # Show regression summary
model.summary()
```

Out[17]:

OLS Regression Results

Dep. Variable:	num_trips	R-squared:	0.687
Model:	OLS	Adj. R-squared:	0.684
Method:	Least Squares	F-statistic:	197.9
Date:	Fri, 18 Apr 2025	Prob (F-statistic):	1.59e-89
Time:	13:39:26	Log-Likelihood:	-3775.7
No. Observations:	365	AIC:	7561.
Df Residuals:	360	BIC:	7581.
Df Model:	4		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
const	2.186e+04	1778.374	12.290	0.000	1.84e+04	2.54e+04
temperature_avg	950.7864	42.906	22.160	0.000	866.409	1035.164
precipitation_sum	-554.2002	72.091	-7.688	0.000	-695.972	-412.428
windspeed_10m_max	-266.0749	81.373	-3.270	0.001	-426.102	-106.048
sunshine_duration	-0.0039	0.034	-0.114	0.909	-0.070	0.062

Omnibus:	13.958	Durbin-Watson:	0.629
Prob(Omnibus):	0.001	Jarque-Bera (JB):	14.860
Skew:	0.494	Prob(JB):	0.000593
Kurtosis:	2.989	Cond. No.	1.52e+05

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

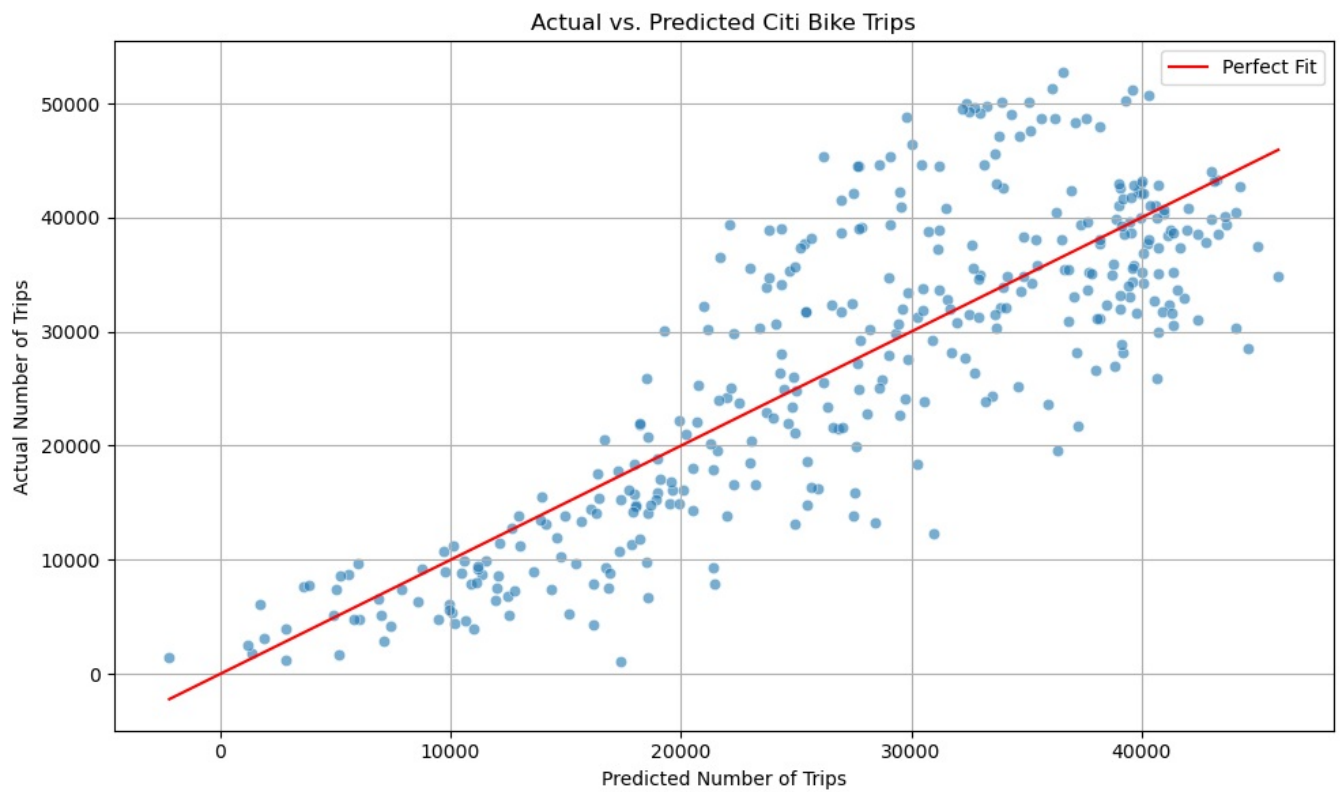
[2] The condition number is large, 1.52e+05. This might indicate that there are strong multicollinearity or other numerical problems.

Visualization

```
In [18]: # Predict values using the fitted model
merged_df['predicted_trips'] = model.predict(X)
```

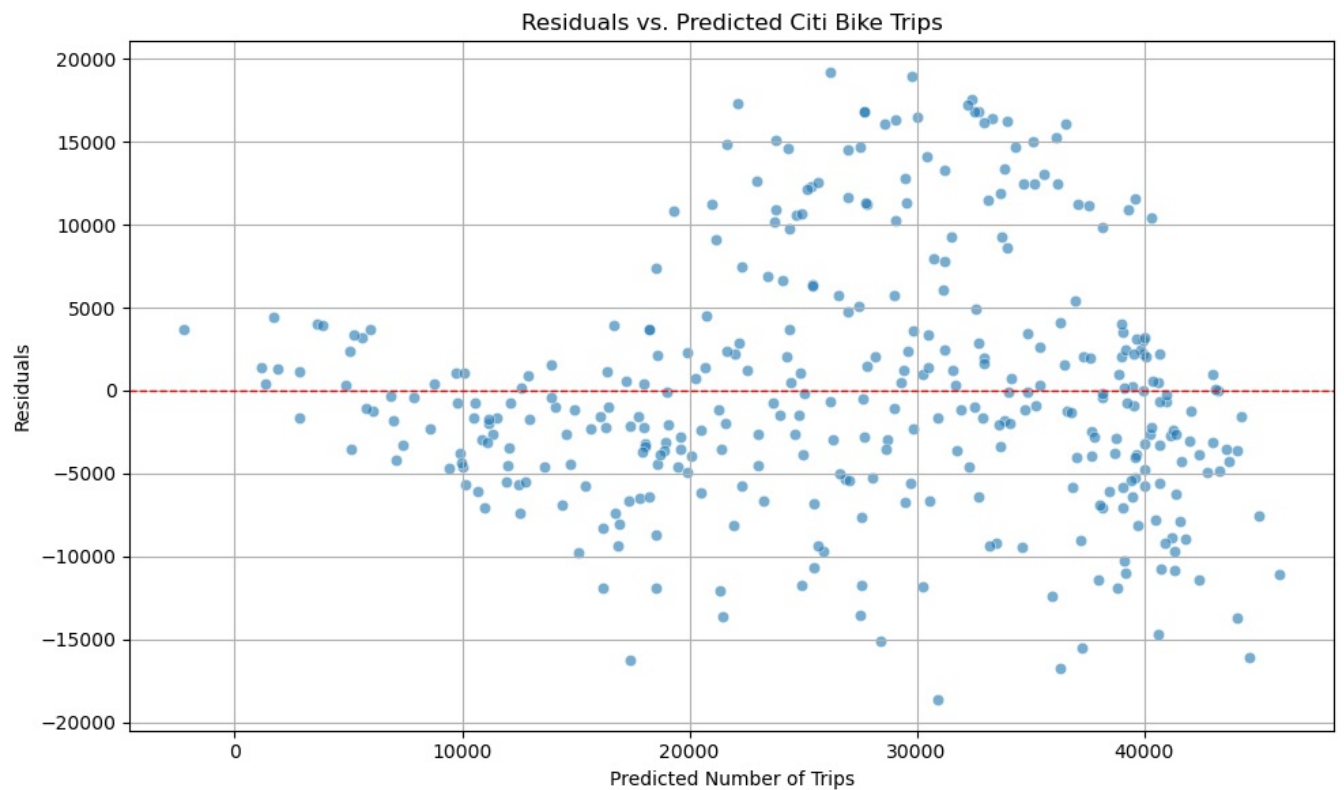
```
In [19]: # Create a new DataFrame just for plotting
plot_df = merged_df[['num_trips', 'predicted_trips']].copy()

# Plot actual vs. predicted values with a perfect fit line
plt.figure(figsize=(10, 6))
sns.scatterplot(data=plot_df, x='predicted_trips', y='num_trips', alpha=0.6)
sns.lineplot(data=plot_df.sort_values('predicted_trips'), x='predicted_trips', y='predicted_trips', color='red')
plt.title('Actual vs. Predicted Citi Bike Trips')
plt.xlabel('Predicted Number of Trips')
plt.ylabel('Actual Number of Trips')
plt.legend()
plt.grid(True)
plt.tight_layout()
plt.show()
```

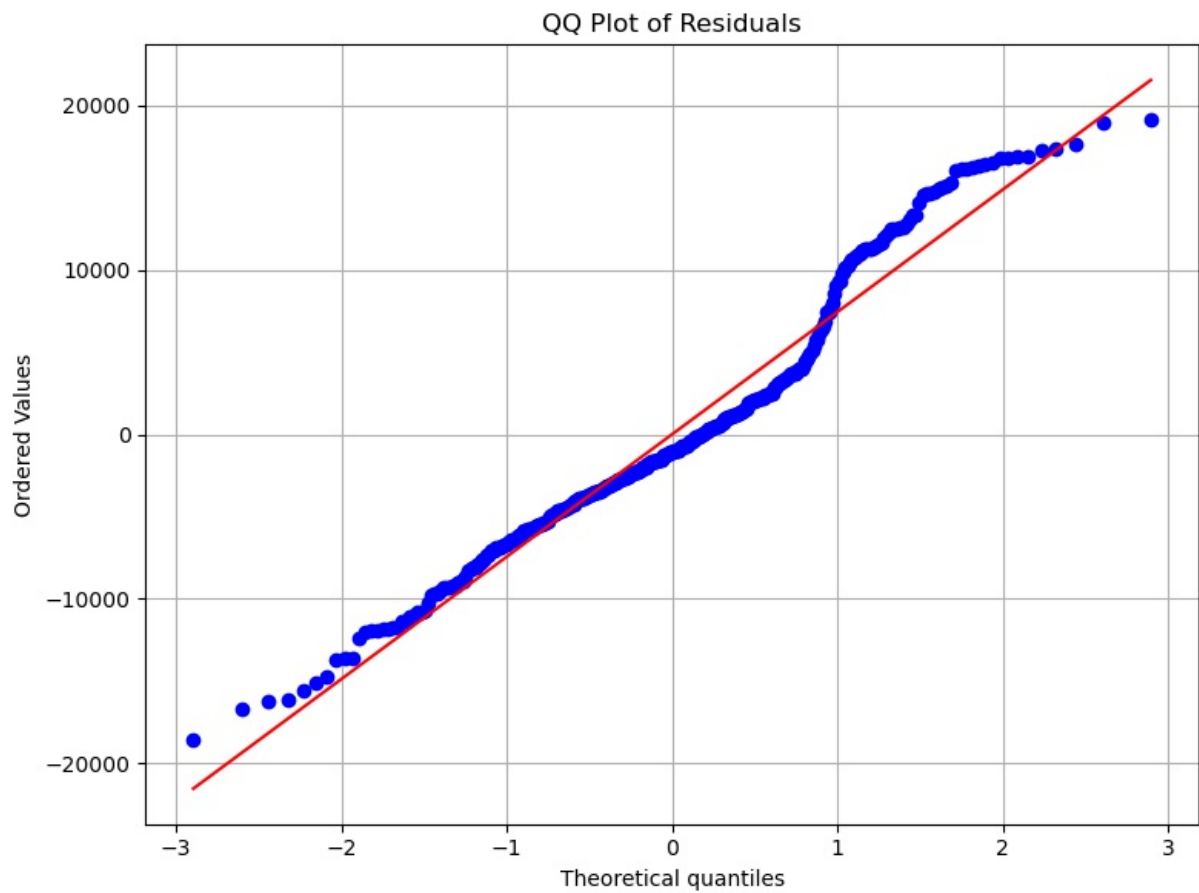


```
In [20]: # Calculate residuals
merged_df['residuals'] = merged_df['num_trips'] - merged_df['predicted_trips']
```

```
In [21]: # Plot residuals vs. predicted values
plt.figure(figsize=(10, 6))
sns.scatterplot(x='predicted_trips', y='residuals', data=merged_df, alpha=0.6)
plt.axhline(0, color='red', linestyle='--', linewidth=1)
plt.title('Residuals vs. Predicted Citi Bike Trips')
plt.xlabel('Predicted Number of Trips')
plt.ylabel('Residuals')
plt.grid(True)
plt.tight_layout()
plt.show()
```



```
In [22]: # Generate QQ plot of residuals
plt.figure(figsize=(8, 6))
stats.probplot(merged_df['residuals'], dist="norm", plot=plt)
plt.title('QQ Plot of Residuals')
plt.grid(True)
plt.tight_layout()
plt.show()
```



Log-transforming

```
In [23]: # Add log-transformed features
merged_df['log_num_trips'] = np.log1p(merged_df['num_trips']) # log(1 + trips)
merged_df['log_temperature_avg'] = np.log1p(merged_df['temperature_avg'] - merged_df['temperature_avg'].min() + 1)
merged_df['log_precipitation_sum'] = np.log1p(merged_df['precipitation_sum'])
merged_df['log_windspeed_10m_max'] = np.log1p(merged_df['windspeed_10m_max'])
merged_df['log_sunshine_duration'] = np.log1p(merged_df['sunshine_duration'])
```

```
In [24]: # Define new model predictors and response
X_log = merged_df[['log_temperature_avg', 'log_precipitation_sum', 'log_windspeed_10m_max', 'log_sunshine_duration']]
y_log = merged_df['log_num_trips']
X_log = sm.add_constant(X_log)

# Fit the log-transformed model
model_log = sm.OLS(y_log, X_log).fit()

# Show summary of the log-transformed model
model_log.summary()
```

Out[24]:

OLS Regression Results							
Dep. Variable:	log_num_trips		R-squared:	0.734			
Model:	OLS		Adj. R-squared:	0.731			
Method:	Least Squares		F-statistic:	248.5			
Date:	Fri, 18 Apr 2025		Prob (F-statistic):	3.63e-102			
Time:	13:39:30		Log-Likelihood:	-162.49			
No. Observations:	365		AIC:	335.0			
Df Residuals:	360		BIC:	354.5			
Df Model:	4						
Covariance Type:	nonrobust						
		coef	std err	t	P> t	[0.025	0.975]
	const	6.7365	0.272	24.805	0.000	6.202	7.271
	log_temperature_avg	1.0576	0.039	27.011	0.000	0.981	1.135
	log_precipitation_sum	-0.1957	0.024	-8.024	0.000	-0.244	-0.148
	log_windspeed_10m_max	-0.1194	0.072	-1.666	0.097	-0.260	0.022
	log_sunshine_duration	0.0330	0.009	3.794	0.000	0.016	0.050
Omnibus:	137.200	Durbin-Watson:	1.050				
Prob(Omnibus):	0.000	Jarque-Bera (JB):	1162.070				
Skew:	-1.334	Prob(JB):	4.57e-253				
Kurtosis:	11.324	Cond. No.	151.				


Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

Visualization


```
In [25]: # Predict log(num trips) and back-transform to original scale
merged_df['log_predicted_trips'] = model_log.predict(X_log)
merged_df['backtransformed_predicted_trips'] = np.expml(merged_df['log_predicted_trips'])
```

```
In [26]: # Plot actual vs. predicted (original scale)
plt.figure(figsize=(10, 6))
sns.scatterplot(
    x='backtransformed_predicted_trips',
    y='num_trips',
    data=merged_df,
    alpha=0.6
)
sns.lineplot(
    x='backtransformed_predicted_trips',
    y='backtransformed_predicted_trips',
    data=merged_df.sort_values('backtransformed_predicted_trips'),
    color='red',
    label='Perfect Fit'
)
plt.title('Actual vs. Predicted Citi Bike Trips (Log-Transformed Model)')
plt.xlabel('Predicted Number of Trips')
plt.ylabel('Actual Number of Trips')
plt.legend()
plt.grid(True)
plt.tight_layout()
plt.show()
```



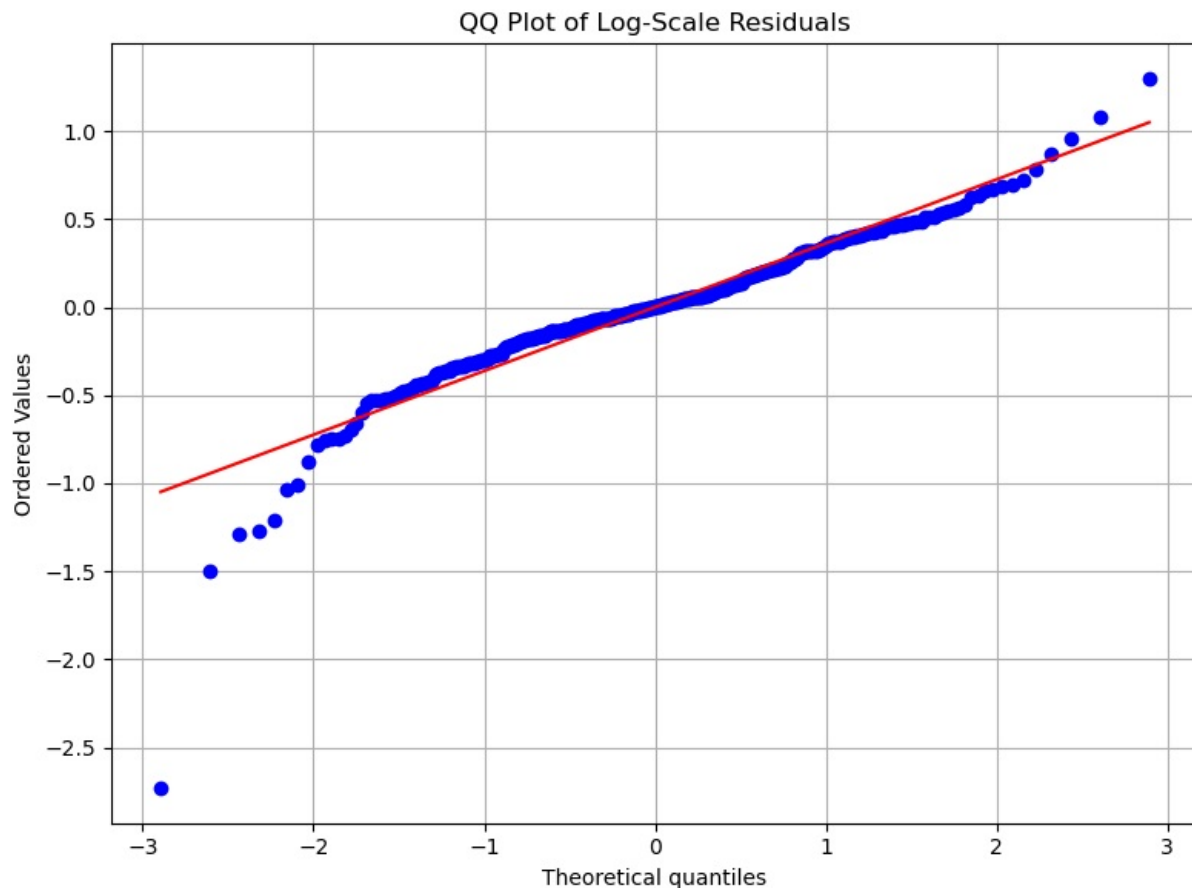
```
In [27]: # Calculate residuals in log scale
merged_df['log_residuals'] = merged_df['log_num_trips'] - merged_df['log_predicted_trips']
```

```
In [28]: # Residual plot: log residuals vs. predicted (log scale)
plt.figure(figsize=(10, 6))
sns.scatterplot(
    x='log_predicted_trips',
    y='log_residuals',
    data=merged_df,
    alpha=0.6
)
plt.axhline(0, color='red', linestyle='--', linewidth=1)
plt.title('Residuals vs. Predicted (Log Scale)')
plt.xlabel('Predicted log(Number of Trips)')
plt.ylabel('Residuals (log scale)')
plt.grid(True)
plt.tight_layout()
plt.show()
```



```
In [29]: # QQ plot for log residuals
plt.figure(figsize=(8, 6))
```

```
stats.probplot(merged_df['log_residuals'], dist="norm", plot=plt)
plt.title('QQ Plot of Log-Scale Residuals')
plt.grid(True)
plt.tight_layout()
plt.show()
```



Partial Least Squares Regression

```
In [10]: merged_df_pp = pd.merge(citibike_df, weather_df, on='date', how='inner').dropna()
```

```
In [11]: features = [
    'temperature_2m_max', 'temperature_2m_min', 'apparent_temperature_max', 'apparent_temperature_min',
    'precipitation_sum', 'rain_sum', 'snowfall_sum', 'precipitation_hours',
    'sunshine_duration', 'windspeed_10m_max', 'windgusts_10m_max', 'winddirection_10m_dominant',
    'shortwave_radiation_sum', 'et0_fao_evapotranspiration', 'temperature_avg'
]
X = merged_df_pp[features]
y = merged_df_pp['num_trips']
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)
X_train_scaled, X_test_scaled, y_train, y_test = train_test_split(X_scaled, y, test_size = 0.2, random_state= 4)
```

```
In [32]: pls = PLSRegression(n_components=5)
pls.fit(X_train_scaled, y_train)

y_pred = pls.predict(X_test_scaled)

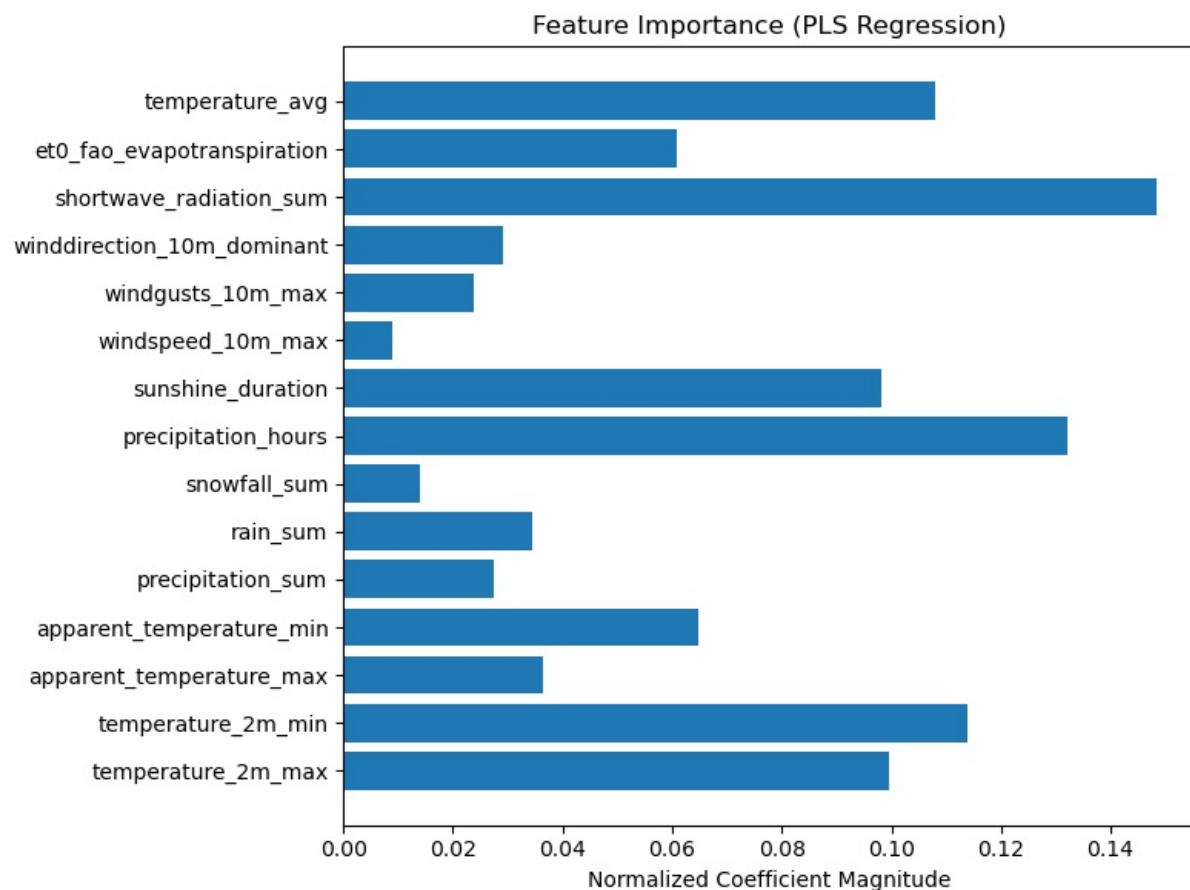
rmse = np.sqrt(mean_squared_error(y_test, y_pred))
r2 = r2_score(y_test, y_pred)

print("RMSE:", rmse)
print("R²:", r2)

RMSE: 6759.398624717034
R²: 0.745126140393312
```


```
In [33]: coeffs = np.abs(pls.coef_).flatten()
coeffs /= np.sum(coeffs)

plt.figure(figsize=(8, 6))
plt.barh(X.columns, coeffs)
plt.xlabel("Normalized Coefficient Magnitude")
plt.title("Feature Importance (PLS Regression)")
plt.tight_layout()
plt.show()
```

```
In [34]: coeffs_signed = pls.coef_.flatten()
impact_df = pd.DataFrame({
    'Feature': features,
    'Coefficient': coeffs_signed
}).sort_values(by='Coefficient', ascending=True)

plt.figure(figsize=(8, 6))
bars = plt.barh(impact_df['Feature'], impact_df['Coefficient'],
                color=['green' if c > 0 else 'red' for c in impact_df['Coefficient']])
plt.axvline(0, color='black', linewidth=0.8)
plt.xlabel("PLS Coefficient (Signed)")
plt.title("Feature Impact on Number of Trips (PLS Regression)")
plt.tight_layout()
plt.show()
```



Tree-Based Model & Feature Importance Analysis

```
In [17]: def evaluate_model(model, X_test, y_test, title="Model"):
y_pred = model.predict(X_test)
rmse = np.sqrt(mean_squared_error(y_test, y_pred))
r2 = r2_score(y_test, y_pred)

print(f"{title} RMSE: {rmse:.4f}")
print(f"{title} R²: {r2:.4f}")

# Plot actual vs. predicted
plt.figure(figsize=(6, 5))
sns.scatterplot(x=y_test, y=y_pred)
plt.plot([y_test.min(), y_test.max()], [y_test.min(), y_test.max()], 'r--')
plt.xlabel("Actual")
plt.ylabel("Predicted")
plt.title(f"{title} - Actual vs. Predicted")
plt.grid(True)
plt.show()

return y_pred

def plot_feature_importance(model, feature_names, title="Feature Importance"):
importances = model.feature_importances_
indices = np.argsort(importances)[::-1]

plt.figure(figsize=(8, 6))
sns.barplot(x=importances[indices], y=np.array(feature_names)[indices])
plt.title(title)
plt.xlabel("Importance")
plt.ylabel("Feature")
plt.grid(True)
plt.show()
```

Decision Tree

```
In [18]: dt_param_dist = {
'max_depth': randint(2, 20),
'min_samples_split': randint(2, 20),
'min_samples_leaf': randint(1, 10)
}

dt = DecisionTreeRegressor(random_state=42)
dt_search = RandomizedSearchCV(dt, dt_param_dist, n_iter=50, cv=5, scoring='r2', n_jobs=-1, random_state=42)
dt_search.fit(X_train_scaled, y_train)
```



```
print("Best Decision Tree Params:", dt_search.best_params_)
best_dt = dt_search.best_estimator_
```

Best Decision Tree Params: {'max_depth': 6, 'min_samples_leaf': 7, 'min_samples_split': 8}

```
In [19]: evaluate_model(best_dt, X_test_scaled, y_test, "Decision Tree")
plot_feature_importance(best_dt, features, "Decision Tree Feature Importance")
```

Decision Tree RMSE: 8008.7662

Decision Tree R²: 0.6422



Random Forest

```
In [21]: rf_param_dist = {
    'n_estimators': randint(50, 200),
    'max_depth': randint(5, 30),
    'min_samples_split': randint(2, 10),
    'min_samples_leaf': randint(1, 10),
    'max_features': ['sqrt', 'log2']
}

rf = RandomForestRegressor(random_state=42)
rf_search = RandomizedSearchCV(rf, rf_param_dist, n_iter=50, cv=5, scoring='r2', n_jobs=-1, random_state=42)
rf_search.fit(X_train_scaled, y_train)
```

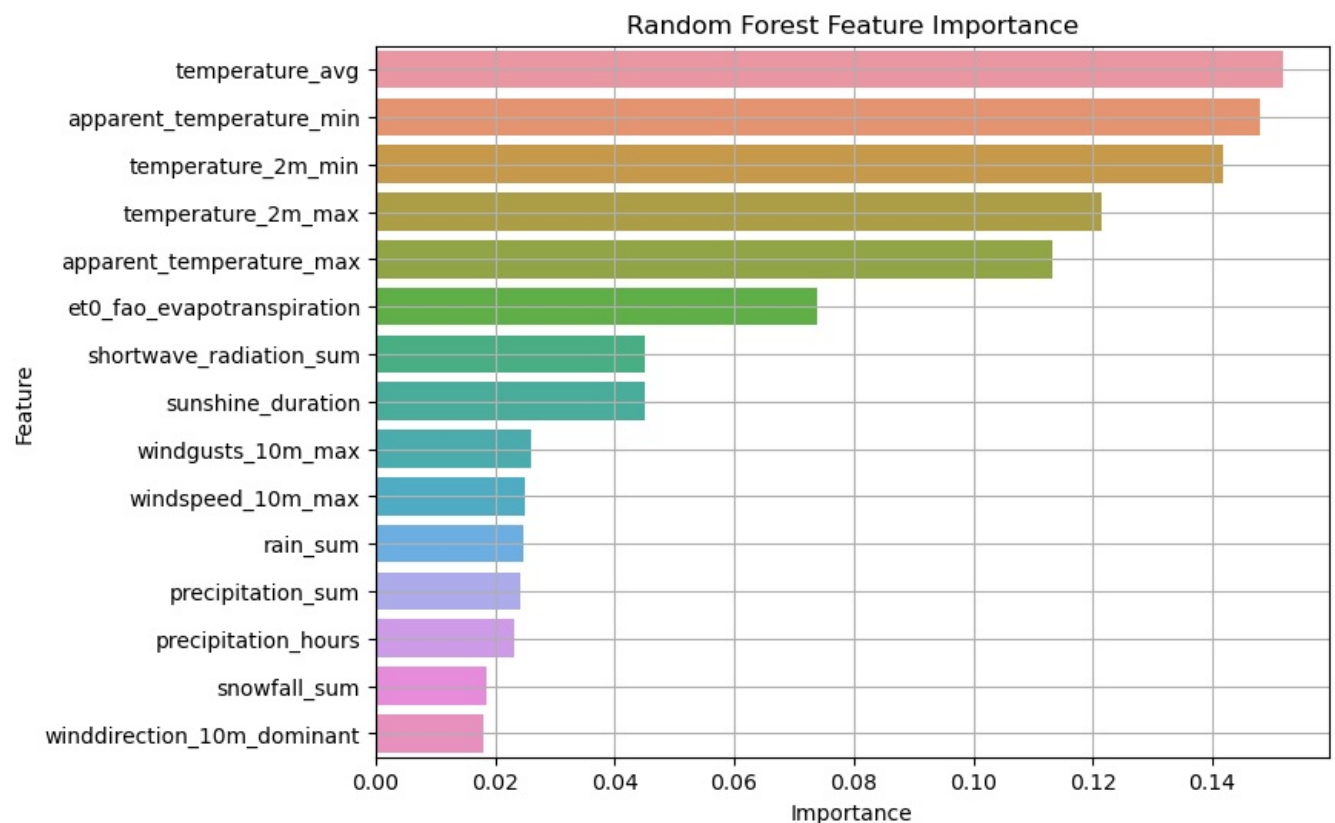

```
print("Best Random Forest Params:", rf_search.best_params_)
best_rf = rf_search.best_estimator_
```

Best Random Forest Params: {'max_depth': 10, 'max_features': 'log2', 'min_samples_leaf': 1, 'min_samples_split': 2, 'n_estimators': 173}

```
In [22]: evaluate_model(best_rf, X_test_scaled, y_test, "Random Forest")
plot_feature_importance(best_rf, features, "Random Forest Feature Importance")
```

Random Forest RMSE: 6682.8512

Random Forest R²: 0.7509



XGBoost

```
In [23]: xgb_param_dist = {
    'n_estimators': randint(50, 200),
    'max_depth': randint(3, 15),
    'learning_rate': uniform(0.01, 0.3),
    'subsample': uniform(0.5, 0.5),
    'colsample_bytree': uniform(0.5, 0.5),
    'gamma': uniform(0, 5)
}
```

```
xgb = XGBRegressor(random_state=42, verbosity=0)
xgb_search = RandomizedSearchCV(xgb, xgb_param_dist, n_iter=50, cv=5, scoring='r2', n_jobs=-1, random_state=42)
xgb_search.fit(X_train_scaled, y_train)
```



```
print("Best XGBoost Params:", xgb_search.best_params_)
best_xgb = xgb_search.best_estimator_
```

Best XGBoost Params: {'colsample_bytree': 0.5384899549143964, 'gamma': 1.4487572645688402, 'learning_rate': 0.058366386176201324, 'max_depth': 4, 'n_estimators': 94, 'subsample': 0.647816842918857}

```
In [24]: evaluate_model(best_xgb, X_test_scaled, y_test, "XGBoost")
plot_feature_importance(best_xgb, features, "XGBoost Feature Importance")
```

XGBoost RMSE: 6847.1848

XGBoost R²: 0.7385



```
In [26]: dt_importance = pd.Series(best_dt.feature_importances_, index=features, name='DecisionTree')
rf_importance = pd.Series(best_rf.feature_importances_, index=features, name='RandomForest')
xgb_importance = pd.Series(best_xgb.feature_importances_, index=features, name='XGBoost')
```

```
# Combine into a single DataFrame
importance_df = pd.concat([dt_importance, rf_importance, xgb_importance], axis=1)
```

```
# Option A: Normalize
normalized_df = importance_df.div(importance_df.sum(axis=0), axis=1)
```

```
# Option B: Rank
rank_df = importance_df.rank(ascending=False)
```

```
# Aggregated importance
normalized_df['Mean'] = normalized_df.mean(axis=1)
rank_df['MeanRank'] = rank_df.mean(axis=1)
```

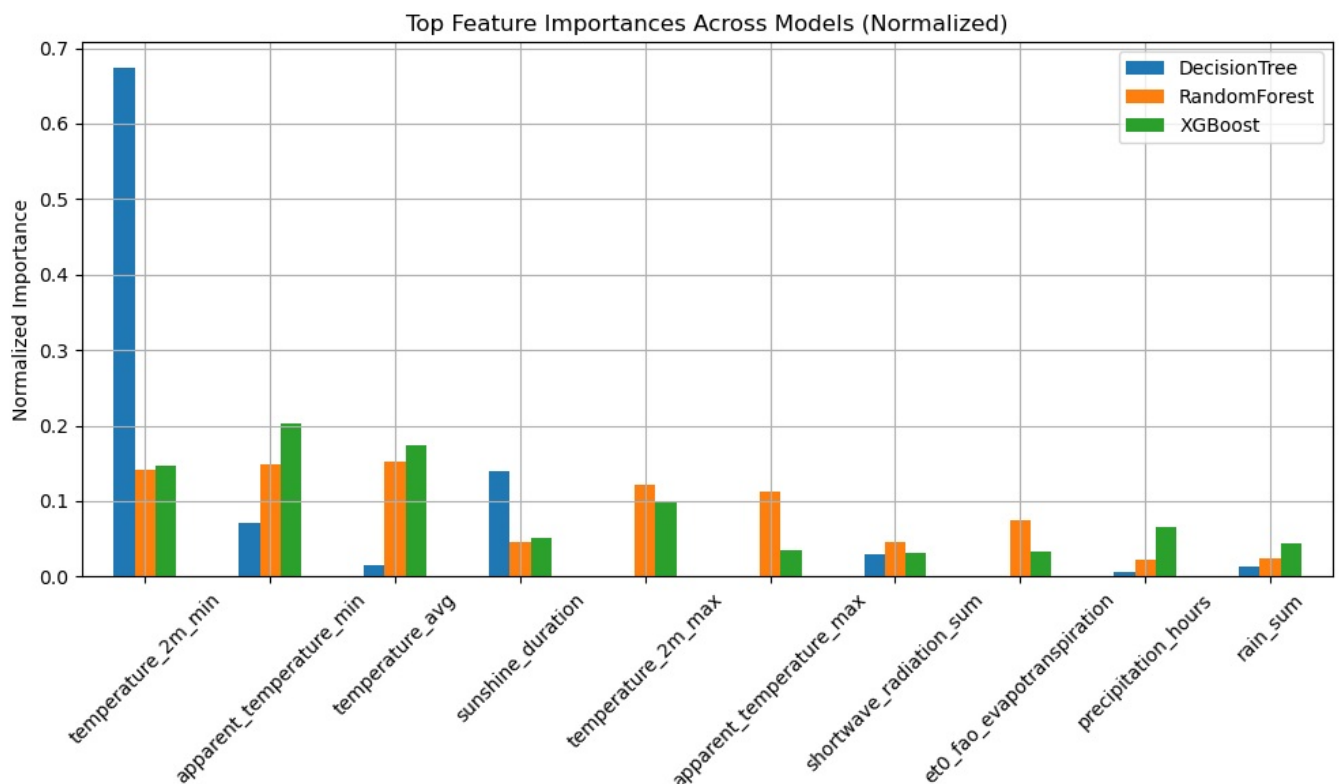
```
# Sort by importance or rank
top_features_by_importance = normalized_df.sort_values('Mean', ascending=False)
top_features_by_rank = rank_df.sort_values('MeanRank')
```

```
# Choose top 5
selected_features = top_features_by_importance.head(5).index.tolist()
print("Selected Top Features (by mean importance):", selected_features)
```

Selected Top Features (by mean importance): ['temperature_2m_min', 'apparent_temperature_min', 'temperature_avg', 'sunshine_duration', 'temperature_2m_max']

```
In [27]: top_n = 10
top_plot_df = normalized_df.sort_values('Mean', ascending=False).head(top_n)

top_plot_df[top_plot_df.columns[:-1]].plot(kind='bar', figsize=(10, 6))
plt.title("Top Feature Importances Across Models (Normalized)")
plt.ylabel("Normalized Importance")
plt.xticks(rotation=45)
plt.grid(True)
plt.tight_layout()
plt.show()
```



Statistical Testing

```
In [ ]: import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from scipy.stats import f_oneway
from statsmodels.stats.multicomp import pairwise_tukeyhsd

# Load data
weather_df = pd.read_csv("nyc_weather_2013_2023.csv")
bike_df = pd.read_csv("citibike_day.csv")

# Convert date columns
weather_df["time"] = pd.to_datetime(weather_df["time"])
bike_df["date"] = pd.to_datetime(bike_df["date"])

# Merge
merged_df = pd.merge(bike_df, weather_df, left_on="date", right_on="time")

# Temperature category
def categorize_temp(temp):
    if temp < 5:
        return "Cold"
```

```

elif 5 <= temp <= 20:
    return "Mild"
else:
    return "Hot"

merged_df["temp_category"] = merged_df["temperature_2m_max"].apply(categorize_temp)

# ANOVA
temp_groups = merged_df.groupby("temp_category")["num_trips"].apply(list)
temp_anova = f_oneway(*temp_groups)
print("Temperature ANOVA:", temp_anova)

# Tukey HSD for Temperature
tukey_temp = pairwise_tukeyhsd(
    endog=merged_df["num_trips"],
    groups=merged_df["temp_category"],
    alpha=0.05
)
print(tukey_temp)

# Boxplot
plt.figure(figsize=(8, 6))
sns.boxplot(x="temp_category", y="num_trips", data=merged_df)
plt.title("Citi Bike Trips by Temperature Category")
plt.xlabel("Temperature Category")
plt.ylabel("Number of Trips")
plt.tight_layout()
plt.show()

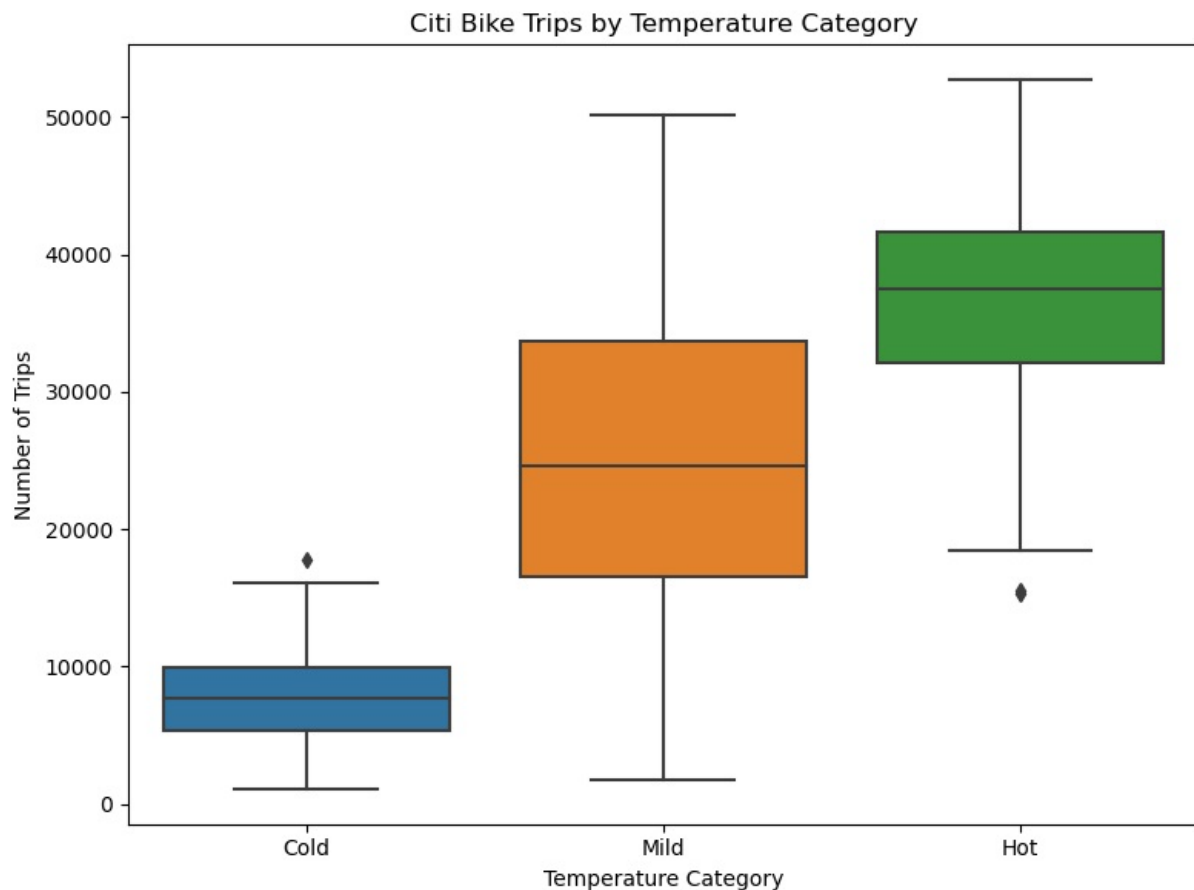
```

Temperature ANOVA: F_onewayResult(statistic=248.93908429347974, pvalue=9.843923705640121e-69)
Multiple Comparison of Means - Tukey HSD, FWER=0.05

```

=====
group1 group2  meandiff  p-adj   lower   upper  reject
-----
Cold   Hot    28598.0609   0.0   25560.4613  31635.6605   True
Cold   Mild   17507.4512   0.0   14447.9848  20566.9176   True
Hot    Mild  -11090.6097   0.0  -13476.9897  -8704.2297   True
=====

```



```

In [38]: # Precipitation category
def categorize_precip(p):
    if p == 0:
        return "No Rain"
    elif p <= 5:
        return "Light Rain"
    else:
        return "Heavy Rain"

merged_df["precip_category"] = merged_df["precipitation_sum"].apply(categorize_precip)

# ANOVA for Precipitation
precip_groups = merged_df.groupby("precip_category")["num_trips"].apply(list)

```


```
precip_anova = f_oneway(*precip_groups)
print("Precipitation ANOVA:", precip_anova)
```

```
# Tukey HSD for Precipitation
tukey_precip = pairwise_tukeyhsd(
    endog=merged_df["num_trips"],
    groups=merged_df["precip_category"],
    alpha=0.05
)
print(tukey_precip)

# Boxplot
plt.figure(figsize=(8, 5))
sns.boxplot(x="precip_category", y="num_trips", data=merged_df)
plt.title("Bike Trips by Precipitation Category")
plt.tight_layout()
plt.show()
```

Precipitation ANOVA: F_onewayResult(statistic=15.62696766833985, pvalue=3.093315843297417e-07)
Multiple Comparison of Means - Tukey HSD, FWER=0.05

group1	group2	meandiff	p-adj	lower	upper	reject
Heavy Rain	Light Rain	6250.2382	0.0125	1101.9523	11398.5242	True
Heavy Rain	No Rain	10601.1689	0.0	5987.9472	15214.3906	True
Light Rain	No Rain	4350.9307	0.0176	614.8757	8086.9856	True



```
In [39]: # Sunshine Duration category
def categorize_sunshine(sun):
    if sun < 10000:
        return "Low Sunshine"
    elif 10000 <= sun < 25000:
        return "Moderate Sunshine"
    else:
        return "High Sunshine"

merged_df["sunshine_category"] = merged_df["sunshine_duration"].apply(categorize_sunshine)

# ANOVA for Sunshine
sunshine_groups = merged_df.groupby("sunshine_category")["num_trips"].apply(list)
sunshine_anova = f_oneway(*sunshine_groups)
print("Sunshine Duration ANOVA:", sunshine_anova)

# Tukey HSD for Sunshine Duration
tukey_sun = pairwise_tukeyhsd(
    endog=merged_df["num_trips"],
    groups=merged_df["sunshine_category"],
    alpha=0.05
)
print(tukey_sun)

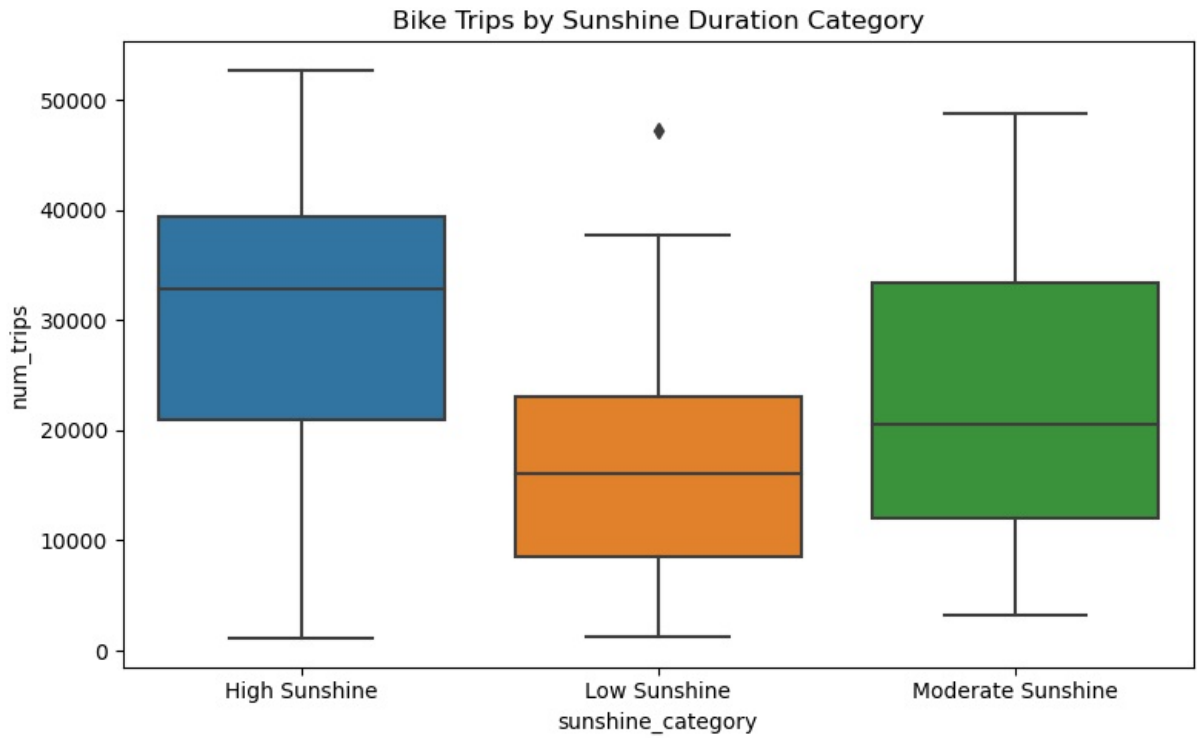
# Boxplot
plt.figure(figsize=(8, 5))
sns.boxplot(x="sunshine_category", y="num_trips", data=merged_df)
plt.title("Bike Trips by Sunshine Duration Category")
plt.tight_layout()
```



```
plt.show()
```

Sunshine Duration ANOVA: F_onewayResult(statistic=25.782967460505027, pvalue=3.401338631787172e-11)
Multiple Comparison of Means - Tukey HSD, FWER=0.05

group1	group2	meandiff	p-adj	lower	upper	reject
High Sunshine	Low Sunshine	-12750.547	0.0	-17123.4046	-8377.6893	True
High Sunshine	Moderate Sunshine	-6559.2737	0.0049	-11447.7829	-1670.7646	True
Low Sunshine	Moderate Sunshine	6191.2733	0.0427	160.6584	12221.8881	True



```
In [ ]:
```